



COMP1002

Data Structures and Algorithms

Lecture 1b: File I/O



Curtin University

Department of Computing

Last updated: [February 20, 2017](#)

Copyright Warning

COMMONWEALTH OF AUSTRALIA

Copyright Regulation 1969

WARNING

This material has been copied and communicated to you by or on behalf of Curtin University of Technology pursuant to Part VB of the Copyright Act 1968 (the Act)

The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice

This Week

- Files and efficient I/O with files
- File streams in Java
- Tokenizers
- Reading and parsing CSV files
- Writing CSV files
- Text vs Binary
- Serialization

Why Files?

- RAM is volatile and private to an application, so it's not a good match for the following purposes:
 - Storing application data long-term (between runs)
 - Sharing information between applications
 - Reading in bulk data provided by a user
- In contrast, files stored on disk are (semi)permanent, can be shared between applications and can be manipulated by the user outside the application

File Input/Output

- Unlike RAM, files are effectively an input to and/or an output from the application
 - Hence the term File I/O, for Input/Output
- Three basic steps in file I/O: (applies to any platform)
 - Open the file
 - Read data from file and/or write data to file
 - Close the file

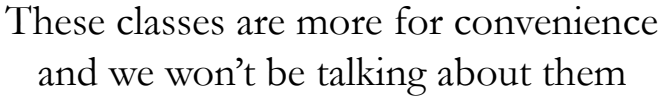
File I/O in Java

- Java encapsulates File I/O into a set of classes
 - All file I/O classes support the three basic steps
 - Different classes are used for manipulating different kinds of files or to perform I/O in different ways
 - » Depends on what you need
 - Other kinds of I/O are also part of the same set of classes
 - » *e.g.*, keyboard input, screen output, network I/O, device I/O, *etc.*
 - » This was done to unify the handling of all the different I/O streams into a single consistent programming interface

Java I/O Classes

- The I/O classes can be split into three groups:
 - Classes that represent an I/O *stream* (e.g., file, device, *etc.*)
 - Classes to read from a stream
 - Classes to write to a stream
- There are specialised versions of the basic classes for each I/O type (file, network, device, *etc.*)
- We will concentrate on the file I/O classes
 - I am just making you aware of the broader scope to help explain why there are so many classes involved!

Java I/O Class Hierarchy



Steps in Java for Reading a File

1. Create a stream object for a file
 - `FileInputStream`
 2. Create an object that can read that stream
 - `InputStreamReader`
 3. Read and process data from the file
 4. Close the `FileInputStream`
- Note that a `FileReader` combines the first two steps into one

Steps in Java for Writing a File

Pretty much the same as for reading

1. Create a stream object for a file
 - `FileOutputStream`
 2. Create an object that can write to that stream
 - `OutputStreamWriter` or `PrintWriter`
 3. Write data to the file
 4. Close the `FileOutputStream`
- Again, `FileWriter` combines the first two steps

Efficient File I/O

- If you look at the methods for `InputStreamReader`, you'll see that it has *two* `read()` methods
 - One for reading a byte at a time
 - » This returns `-1` if no more bytes to read (*i.e.*, end-of-file reached)
 - A second for reading an array of bytes at a time
- Since it's rare that we ever want just a single byte of data, the first method seems less useful
- However, when trying to read in text data there are conceptual difficulties with reading chunks at a time

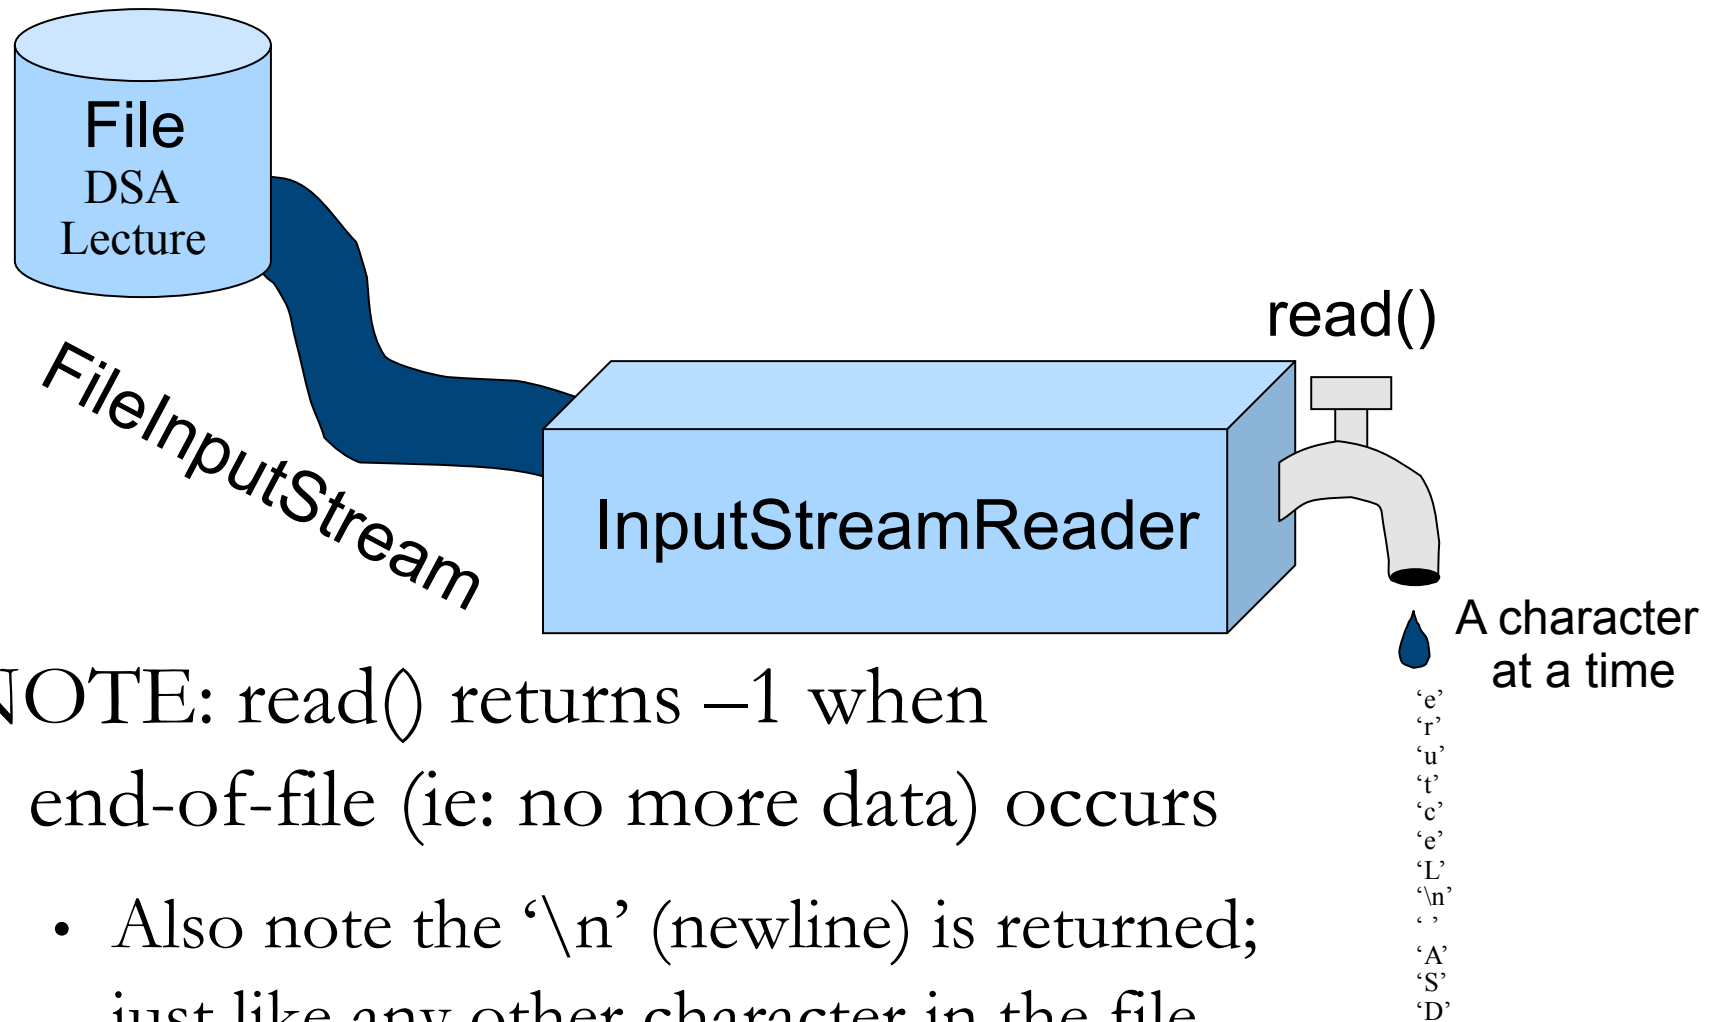
Text vs Binary Data

- Data files generally fall into two broad categories:
 - Files containing text data (relatively unstructured)
 - Files containing **binary** data (everything else)
- Binary data is highly structured
 - *e.g.*, images, databases, executable files, *etc.*
 - There is a lot of prior knowledge on precisely where information exists in the file and how large blocks are
- Text data is *un*structured, so it is usually impossible to know beforehand how many bytes to read
 - *e.g.*, how many characters in an arbitrary line of text?

Efficient File I/O

- So reading blocks of N bytes from a binary file is OK
 - You will know beforehand how many bytes are needed
- But when reading text, you have to constantly be ready for the end of data since you can't predict it
 - *e.g.*, end of word, end of line, or end of file
- One approach is to read data in a byte at a time, and check each byte for end-of-X (' ', '\n' or -1)
 - The problem is that this will be very slow since hard disks are fastest at reading blocks of data at a time
 - » It's like filling a bucket one drop at a time vs opening the tap and letting the water flow freely into the bucket

InputStreamReader




NOTE: `read()` returns `-1` when end-of-file (ie: no more data) occurs

- Also note the `'\n'` (newline) is returned; just like any other character in the file

How Hard Disks Work

- To understand file access efficiency, you have to understand how a disk drive reads and writes data
 - Applies to CDs/DVDs as well, but not to SSDs.
 - » SSD = solid state drive. These behave more like RAM
 - » HDD = hard disk drive
- A hard disk drive is a set of circular disks (“platters”) that can apply and retain magnetism on its surface
 - Data is stored via this magnetism: magnetic = 1, not = 0.
- The magnetism is manipulated by a read-write head
 - An arm with a small magnetometer on its tip

Internals of a Hard Disk



Data Organisation on Hard Disks

- A drive is made up of multiple platters and heads
- Data is organised in concentric circles around the disk, called **tracks** (or alternatively, cylinders)
 - A cylinder is the set of equivalent tracks across all platters
- Each track is split into pieces called **sectors**
- The O/S then maps these tracks and sectors to blocks
 - Blocks are the smallest chunk of data that the O/S will read/write to/from disk. Usually 512 bytes or 1Kb
- Nowadays, the HDD itself maps tracks/sectors to an internal config, but that is not seen by you or the O/S

Hard Disk Access Speed

- Accessing data is a two-stage process
 - First the **head** must be positioned on the track where the data resides (*somewhere* on that track)
 - Then the head must wait until the disk, already spinning, brings the target sector under the head so it can be read
- Usually files are organised so that their data is stored in sequential sectors on the same track(s) (if possible)
- So there are two important elements to HDD speed:
 - **Seek time** – how long the head takes to get to a track
 - **Transfer time** – how fast can the data be streamed off the track in sequential order (faster spin = faster reading)

Modern Access Speeds

- Seek time is worst: on average around 7ms
 - Physical limits of motors/actuators
- Transfer speed increases with HDD density (size)
 - So 50Mb/s is common nowadays
- Thus when dealing with files, you want to minimise the amount of seeking and maximise streaming
 - Not totally in your control – the O/S is the one that decides where files go based on what sectors are available
 - But modern O/Ss do a good job of locating files contiguously to avoid fragmentation of the files

Implications for File I/O Efficiency

- When a programmer reads N bytes from a file, the O/S determines the track/sector of that file and sends a read instruction to the HDD for the N bytes
 - The HDD will then seek and read N bytes in one operation
 - This will (ideally) involve one seek and one stream read
 - » (actually, modern HDDs do all sorts of optimisations, but if the file is contiguous it will almost surely be read in one operation)
 - The read data will be sent to RAM by the HDD
 - The O/S will be alerted to the fact that the data is now available and return it to the program

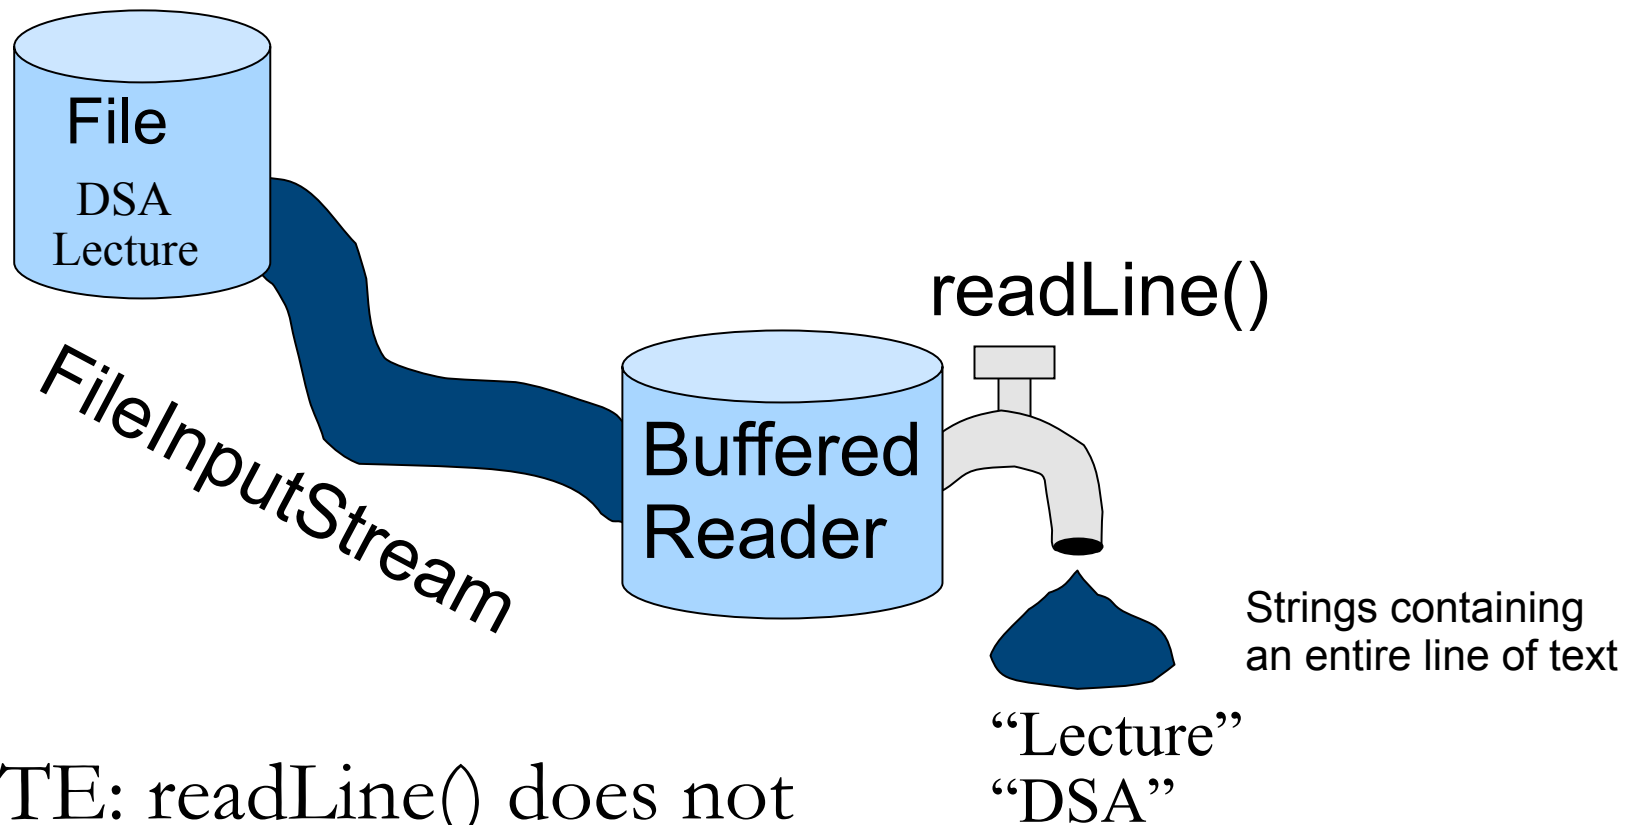
Implications for File I/O Efficiency

- If instead you ask for the N bytes *one byte at a time*, the O/S will submit N read requests to the HDD
 - And each subsequent read will be sent too ‘late’ – the next sector will have spun past the head and the HDD will need to wait until the disk spins back around to read it
 - Worse, there may be another process running that asks for data from a different part of the disk
 - » This will be interleaved with your N reads
 - » It will cause the head to flit back and forth across the disk to access the two different files, racking up the seek time overheads
- So reading a file a single byte at a time will be *slow*

BufferedReader

- Reading a file a single byte at a time is very slow.
- BufferedReader reads in chunks of data, *buffering* it in memory to search for end-of-X
 - » In particular, the `readLine()` method looks for end-of-line
- The idea is to read in ('buffer') a chunk of (say) 1024 bytes, then search these bytes in RAM one at a time for the '`\n`'
 - » `\n` = newline character, marking the end of a line of text
 - » Since RAM is much faster than disk, this is much more efficient
- When a line is found it is extracted and returned, but the rest of the buffer is *kept in memory*
 - » Then for subsequent `readLine()` calls, it first checks the buffer for the '`\n`' before reading another 1024-byte chunk from the file
- Note: you create a `BufferedReader` from another `Reader`

BufferedReader



NOTE: `readLine()` does not include the `'\n'` in the returned lines and returns **null** when end-of-file occurs

File Reading: One Line at a Time

```
METHOD readFileExample IMPORT filename EXPORT NOTHING
```

```
theFile ← OPENFILE filename
```

```
lineNum ← 0
```

```
INPUT line FROM theFile
```

```
WHILE NOT (theFile = EOF)
```

```
    lineNum ← lineNum + 1
```

```
    processLine(line)
```

```
    INPUT line FROM theFile
```

```
ENDWHILE
```

```
CLOSEFILE theFile
```

```
ENDMETHOD
```

← EOF = end of file. Detecting this is *very* language-specific

← Whatever processing on the line is required

← Read the next line

← Close the file

File Reading: One Line at a Time

```
private void readFileExample(String inFilename) {
    FileInputStream fileStrm = null;
    InputStreamReader rdr;
    BufferedReader bufRdr;
    int lineNum;
    String line;
    try {
        fileStrm = new FileInputStream(inFilename);
        rdr = new InputStreamReader(fileStrm);
        bufRdr = new BufferedReader(rdr);

        lineNum = 0;
        line = bufRdr.readLine();
        while (line != null) {
            lineNum++;
            processLine(line);
            line = bufRdr.readLine();
        }
        fileStrm.close();

    }
    catch (IOException e) {
        if (fileStrm != null) {
            try { fileStrm.close(); } catch (IOException ex2) { } // We can't do anything more!
        }
        System.out.println("Error in file processing: " + e.getMessage());
    }
}
```

← Open the file
← Create a reader to read the stream
← To read the stream one line at a time

← Read the first line
← While not end-of-file, process and read lines

← Whatever processing on the line is required
← Read the next line

← Clean up the stream

← MUST catch IOExceptions
← Clean up the stream if it was opened

← Or do a throw

Notes on File Handling

- Make sure you close the file as soon as possible
 - The operating system must track what files are open
 - » The O/S remembers where you were in the file, etc
 - The resources available for this tracking are limited
 - » Run out and the O/S will terminate your program
 - Thus don't leave files open – clean them up *early*
 - » Java doesn't free objects immediately – it waits for the garbage collector, so *always* explicitly `close()` a file once finished with it

Notes on File Handling

- You *must* handle IOException
 - IOExceptions must be caught
 - » Checked Exception
 - » The compiler will complain otherwise
 - » This also forces us to do the try..catch around close()
 - Note that you can add a `throw IOException` clause to the method to ‘avoid’ having to catch IOExceptions
 - » It just means that now the *calling* method must catch them

Parsing

- When dealing with text, it's often necessary to take it apart and organize it ready for processing/storage
- This is called **parsing** – to determine and extract the structure of a piece of text
 - The word originally comes from syntax analysis of written language
- Examples of where parsing is needed:
 - Natural language processing (*e.g.*, spelling/grammar checks)
 - Building Web search indexes
 - Compilers: must parse code to detect stmts and variables

Tokenizing

- Tokenizing is the first step in parsing: the process of breaking up a stream of text into basic elements
 - » We'll use US spelling (with a 'z') to avoid confusion
- These elements are called tokens, and what they are depends on what the application is parsing
 - » *e.g.*, Single words, entire lines, equation terms, *etc.*
- Tokens are broken up by searching for character(s) that *delimit* the boundary of a token
 - » *e.g.*, lines are separated by a '\n' newline character
 - » *e.g.*, words are separated by spaces, commas and periods
 - » *e.g.*, equation operands are separated by operators +/*-

Tokenizing with Java

- Java provides two classes to assist in tokenizing
 - StringTokenizer
 - StreamTokenizer
 - They do a similar job except one works on Strings while the other works on Readers
 - » We'll focus on StringTokenizer
- Java also provides the `split(String regex)` method in the `String` class for simple tokenizing of a `String`
 - It's easier to use than the tokenizers, but won't return the delimiting characters (which can be needed)
 - and you don't know what a regex is!

Comma Separated Values

- We'll take a little detour and introduce comma separated values
 - We'll use this as an example for StringTokenizer later on
- We often need to store data to a file
 - The question is, in what form should we store it?
- If the data is in table or matrix form, one can write it out as a set of rows and columns in a certain format:
 - One row is written per line
 - Each row contains multiple fields, one per column
 - Each field's value is separated by a comma ','

CSV Example

Content of File

The diagram illustrates the structure of a CSV file. It features a table with 8 rows and 4 columns. The first row is the header, and the subsequent 7 rows contain data. Callouts identify the components: 'Title' points to the header row; 'Label Serie 1', 'Label Serie 2', and 'Label Serie 3' point to the columns 'Africa', 'Asia', and 'Europe' respectively; 'X-values' points to the first column of data (months); 'Values Serie 1', 'Values Serie 2', and 'Values Serie 3' point to the columns of numerical data.

Sales per region	Africa	Asia	Europe
Jan	34	67	56
Feb	36	87	78
Mar	31	56	88
Apr	29	67	92
Mar	43	56	78
May	54	71	68
Jun	42	65	82

```
Sales per region,Africa,Asia,Europe
Jan,34,67,56
Feb,36,87,78
Mar,31,56,88
Apr,29,67,92
Mar,43,56,78
May,54,71,68
Jun,42,65,82
```

The CSV File

CSV Notes

- The delimiting commas have no trailing space
 - It might look better to have a space after the comma, but it just makes it harder to parse when reading the file later!
- Column sizes don't have to be consistent across rows
 - *e.g.*, the first row (headings) has much longer fields than the same columns in subsequent rows
 - » It would be a waste of space to pad out the fields
 - » Parsing/tokenizing must handle these variable-length fields
- Numeric data is converted into its textual equivalent
 - If we saved integers directly, we might get things that look like text
'\n' (ASCII 13) or ',' (ASCII 44) but are merely part of the data

Text vs Binary

- You don't *have* to save table data in CSV format
 - In fact, dumping raw binary data is often more efficient, and you know how large each field is (*e.g.*, ints = 4 bytes)
- However, CSV text data has some advantages:
 - Easy for humans to read and edit
 - Highly portable to different platforms
 - » big endian vs little endian
 - Fields are explicitly separated with commas
 - CSV is a widely-known format
 - » XML is now beginning to supplant CSV as the standard format for data interchange between companies, but CSV is still used

StringTokenizer and CSV Data

- So let's parse a CSV file with a StringTokenizer
 - StringTokenizer works a lot like an iterator (unfortunately, it's from JDK 1.0 and so isn't an actual Iterator – it just has similar usage)
 - The idea is to create a StringTokenizer based on the string to parse
 - Then repeatedly ask the StringTokenizer for the next token, until no more tokens exist
 - » And for each token, perform appropriate processing
 - Note that we need to create a new StringTokenizer for every line that is read from the CSV file
 - » So we'll put the tokenizing code into the processLine() method

Parsing a Single CSV Row

```
private void processLine(String csvRow) {  
    String thisToken = null;  
    StringTokenizer strTok;
```

```
    strTok = new StringTokenizer(csvRow, ",");  
    while (strTok.hasMoreTokens()) {  
        thisToken = strTok.nextToken();  
        System.out.print(thisToken + " ");  
    }  
    System.out.println("");  
}
```

← Initialise tokenizer to parse the csvRow

← Iterate over tokens

← Print out each column for this row

Output

- Given the following CSV data file:

```
97452,James,88,96,82,86  
99576,Alan,6,46,34,38  
99888,Geoff,100,68,72,75
```

- The following would be the output of having `readFileExample()` call the shown `processLine()`:

```
97452 James 88 96 82 86  
99576 Alan 6 46 34 38  
99888 Geoff 100 68 72 75
```

Parsing CSV Files with Known Format

- Combine the previous method with the line-reading algorithm for a *general-purpose* CSV parser
- But what if you have very specific data?
 - Consider the following CSV file contents:
97452,James,88,96,82,86
99576,Alan,6,46,34,38
99888,Geoff,100,68,72,75
 - Note that the *meaning* of each column is not in the CSV
 - » We must get this information from somewhere else and hardcode it into the tokenizing/parsing application (clumsy!)
 - » Format for the above data happens to be:
<StdntID>,<Name>,<Assign%>,<Test%>,<Exam%>,<Overall%>

Parsing CSV Files with Known Format

- In this case, we know exactly what each column should contain
- Thus we may want to handle each column specially
 - The next slide considers an application that prints the correct label in front of each field

Parsing a Known-Format CSV Row

```
private void processLine(String csvRow) throws IllegalStateException
{
    int id;
    String name;
    double assign, test, exam, overall;
    StringTokenizer strTok;

    strTok = new StringTokenizer(csvRow);
    try {
        id = Integer.parseInt(strTok.nextToken());
        name = strTok.nextToken();
        assign = Double.valueOf(strTok.nextToken());
        test = Double.valueOf(strTok.nextToken());
        exam = Double.valueOf(strTok.nextToken());
        overall = Double.valueOf(strTok.nextToken());

        System.out.println(" ID:" + id + "Name:" + name + "Assign:" + assign +
                           "Test:" + test + "Exam:" + exam + "Overall:" + overall);
    }
    catch (Exception e) {
        throw new IllegalStateException("CSV row had invalid format");
    }
}
```

← Initialise tokenizer to parse the csvRow

← DON'T loop: just assume the format is exactly what we expect it to be, catching exceptions if it turns out to be a bad format

Output

- Given the CSV data from a previous slide:

97452,James,88,96,82,86

99576,Alan,6,46,34,38

99888,Geoff,100,68,72,75

- The following would be the output of having `readFileExample()` call the last `processLine()`:

ID:97452 Name:James Assign:88 Test:96 Exam:82 Overall:86

ID:99576 Name:Alan Assign:6 Test:46 Exam:34 Overall:38

ID:99888 Name:Geoff Assign:100 Test:68 Exam:72 Overall:75

Writing Text Files

- Writing files is actually simpler than reading them
 - Since you don't have to worry about parsing!
- The overall approach is the same:
 - Open a `FileOutputStream`
 - Create a `Writer`
 - Output data to the file using the `Writer`
- One thing to be careful with is that you must ensure that newlines and commas are put in the right place
 - Assuming you are outputting CSV format, of course!

PrintWriter

- You can use the basic `OutputStreamWriter`, but it's a bit clunky
 - The `write()` method requires an array of bytes, so you have to copy your data into an array before writing
- It would be easier if you could write to a file in the same way that you can print messages to the console
 - Fortunately you can: `PrintWriter`
 - In fact, `System.out` *is* an instance of a `PrintWriter`!
 - Thus writing to file can be identical to printing to screen
 - » Just don't forget to `close()` the file after you are done!!

PrintWriter Example

- The next slide shows an example of writing to a file
- It assumes you have parsed in the student marks info that was read in from the earlier CSV file
 - *e.g.*, we are saving one line of the data to a new CSV file
 - Note that the commas are inserted between the fields, and there are no spaces added

Writing a CSV Row

```
private void writeOneRow(String filename, int ID, String name, double assign, double test,
                        double exam, double overall) {
    FileOutputStream fileStrm = null;
    PrintWriter pw;

    try {
        fileStrm = new FileOutputStream(filename);           ← Open the file for writing
        pw = new PrintWriter(fileStrm);                     ← Initialise writer

        pw.println(id + "," + name + "," + assign + "," + test + "," + exam + "," + overall);
        pw.close();                                         ← Clean up the stream
    }
    catch (IOException e) {                                ← MUST catch IOExceptions
        if (fileStrm != null) {                             ← Clean up the stream if it was opened
            try { fileStrm.close(); } catch (IOException ex2) { } // We can't do anything more!
        }
        System.out.println("Error in writing to file: " + e.getMessage()); ← Or do a throw
    }
}
```

Things Can Be A Bit Simpler

- We have explored the ‘long’ way to open files for I/O
- Java provides other classes and constructors that do some of the steps for you
 - *e.g.*, `FileReader`, `FileWriter`
 - *e.g.*, the `FileOutputStream` creation can be done by an alternate c’tor in `PrintWriter` that accepts the filename
- I don’t use these convenience classes/c’tors so as to explicitly show you how Java is put together
 - But you can use them in your practicals and assignment

Text vs Binary

- Text advantages over binary:
 - Human-readable and editable
 - Text representation of numbers, strings are unambiguous
 - Slight variations in format can often be handled
 - » *e.g.*, if a space is added after the ',' in CSV it could be ignored
- Binary advantages over text:
 - Smaller size for storing the same data
 - » *e.g.*, $\pi = "3.141592653589793"$ in text (17 bytes), but π can be stored with similar precision in only 8 bytes using a double
 - Faster – no need to convert numbers to/from text strings
 - Simpler to parse since exact format and sizing of each piece of data is known in advance

ShippingContainerClass

```
public class ShippingContainerClass
{
    private String name;
    private int serial;
    private char type;
    private double weight;
    private Location location;

    //normal constructors, getters and setters
    //not relevant to this example

    void open(String fileName);
    void save(String fileName);
}
```


Writing Binary Data

– DataOutputStream

- Created from a FileOutputStream
- Includes the following methods
 - » `void writeUTF(String s);`
 - » `void writeInt(int ii);`
 - » `void writeChar(char ch)`
 - » `void writeDouble(double d);`
 - » `void flush();`
 - » `void close();`
 - » and many more

ShippingContainerClass save method

```
private void save(String filename) {
    FileOutputStream fileStrm;
    DataOutputStream dataStrm;

    try {
        fileStrm = new FileOutputStream(filename); ← Underlying stream
        dataStrm = new DataOutputStream(fileStrm); ← Data stream

        dataStrm.writeUTF(name);
        dataStrm.writeInt(serial);
        dataStrm.writeChar(type);
        dataStrm.writeDouble(weight); ← save binary data to filename

        fileStrm.close(); ← Clean up
    }

    catch (IOException e) {
        if (fileStrm != null) { ← Clean up the stream if it was opened
            try { fileStrm.close();
            }
            catch (IOException ex2) { } ← We can't do anything more!
        }
        System.out.println("Error in file processing: " + e.getMessage()); ← Or do a throw
    }
}
```

Reading Binary Data

– DataInputStream

- Created from FileInputStream
- Includes the following methods:
 - » `char readChar();`
 - » `double readDouble();`
 - » `int readInt();`
 - » `string readUTF();`
 - » `void close();`
 - » and many more

ShippingContainerClass open(String fileName)

```
private void open(String filename) {
    FileInputStream fileStrm;
    DataInputStream dataStrm;

    try {
        fileStrm = new FileInputStream(filename);
        dataStrm = new DataInputStream(fileStrm);

        this.name = dataStrm.readUTF();
        this.serial = dataStrm.readInt();
        this.type = dataStrm.readChar();
        this.weight = dataStrm.readDouble();

        fileStrm.close();

    } catch (IOException e) {
        if (fileStrm != null) {
            try { fileStrm.close(); }
            catch (IOException ex2) { }
        }
        System.out.println("Error in file processing: " + e.getMessage());
    }
}
```

← Underlying stream
← Data stream
← read data into classfields
← Clean up
← Clean up the stream if it was opened
← We can't do anything more!
← Or do a throw

Serialization

- Persistence
- Text vs binary for serialization
- Java's object serialization mechanism
- Serialization streams

Persistence

- **Persistence:** The saving and loading application data so that it lasts between executions of the application
 - *i.e.*, the task of storing an application's state to disk and later loading it back to continue running
- We've already had a look at using text and binary formats to save data to disk and read it in later
 - Loading may involve parsing (text) saved data to translate it into a form that is suitable for manipulation by the application
 - Not all persistence requires parsing, but most do simply because storage is static whereas manipulation is dynamic

The Need for Serialization

- What if ContainerClass has a classfield which is an object of another class (aggregation)?
 - it does! Location class
- What if ContainerClass had an inheritance relationship that has classfields in the super class?

Serialization Issues

- Every Java object is a reference (pointer) – even Strings!
- Pointers let you jump around in memory (*i.e.*, non-contiguous)
 - » This is so useful that we simply cannot do without pointers!
- But you cannot save a pointer and expect it to be useful after re-loading it since the pointer says where the data *used to be* in the *old* run of the application!
 - » When the application runs again, it won't place objects in the same place – every time **new** is called, it returns a different address
 - » Furthermore, the object doesn't even exist yet – it is about to be loaded from disk!
- In other words, pointers cannot be made persistent

How Does Serialization Work

- So how could it possibly work?
 - The idea is pack the data to be saved into a single contiguous block so that it can be written all together
 - Java is able to do this *automatically* via a language feature called reflection
- Reflection is the ability for an object to query the methods and fields of another class *at runtime*
 - This returns the names of fields, methods, parameters, return types, etc, and the ability to actually set/get the value of each field
 - » Thus Java can figure out what data fields an object has and pack them into a byte array for saving

Java Serialization

- writing your own serialization code in Java is quite involved
 - The serialization buffer must be an array of bytes
 - Then you must copy every byte of an object into the serialization byte array
- This is why Java's serialization mechanism is so very useful
 - It does all the work of formatting and copying for you
 - From the programmer's viewpoint, it all comes down to the Serializable interface and a couple of Stream classes

Serializable Interface

- To let Java serialize your class, simply make it inherit from the Serializable interface *e.g.*,

```
public class ContainerClass implements Serializable
{
    ...
}
```

- That's it! **You do not have to write any code** – Java will determine the contents of your class and serialize it for you *entirely automatically* using reflection
 - It is possible to override Java's serialization if needed
- Note that Serializable is an *empty* interface! In Java:

```
public interface Serializable { }
```

- *i.e.*, it is only used to mark a class as serializable for Java

Object I/O Streams

- Any class that implements Serializable can then be written and read via two serialization-specific classes
 - **ObjectOutputStream** – for serializing an object and writing it out to an underlying stream
 - **ObjectInputStream** – for reading an underlying stream, deserializing the object and returning it
 - » This creates the object and sets its fields to the serialized data
 - The two classes both know the serialization format
- When creating an Object stream, you must provide another stream – the underlying ‘physical’ stream
 - Usually FileOutputStream and FileInputStream

Object I/O Streams

- ObjectOutputStream will serialize the given object *and all objects referenced by that object*
 - This is done recursively, so that the entire object tree will be serialized
 - This supports the composition relationship: a car object consists of wheel objects, engine object and chassis object
 - » So when serializing a car, we must serialize all contained objects
 - Objects not marked as Serializable won't be serialized
 - » You can also add in custom serialization/deserialization steps, but this is only needed in special circumstances
- ObjectInputStream will deserialize all contained objs

Serializing an Object

```
private void save(ContainerClass objToSave, String filename)
{
    FileOutputStream fileStrm;
    ObjectOutputStream objStrm;

    try {
        fileStrm = new FileOutputStream(filename);           ← Underlying stream
        objStrm = new ObjectOutputStream(fileStrm);         ← Object serialization stream
        objStrm.writeObject(objToSave);                    ← Serialize and save to filename
                                                         This will also save the ContainerClass'
                                                         contained Location object
        objStrm.close();                                    ← Clean up
    }
    catch (Exception e) {                                  ← should do more here
        throw new IllegalArgumentException("Unable to save object to file");
    }
}
```

Deserializing an Object

```
private ContainerClass load(String filename) throws IllegalArgumentException
{
    FileInputStream fileStrm;
    ObjectInputStream objStrm;
    ContainerClass inObj;

    try {
        fileStrm = new FileInputStream(filename);           ← Underlying stream
        objStrm = new ObjectInputStream(fileStrm);         ← Object serialization stream
        inObj = (ContainerClass) objStrm.readObject();      ← Deserialize. Note the cast is needed
        objStrm.close();                                   ← Clean up
    }
    catch (ClassNotFoundException e) {
        System.out.println("Class ContainerClass not found" + e.getMessage);
    }
    catch (Exception e) {
        throw new IllegalArgumentException("Unable to load object from file");
    }
    return inObj;
}
```

XML Serialization

- The Object<Input | Output>Stream classes use a format that is specific to Java
 - Don't expect to be able to pass the object to another language and deserialize it, even if the classes are identical!
 - » In fact, backwards compatibility *in Java* has even been broken
- An alternative that is more portable is XML format
 - XML has become a defacto standard for serializing
 - Java provides XMLEncoder and XMLDecoder for this, which work similarly to the Object<Input | Output>Stream classes
 - JAXB (Java Architecture for XML Binding) is a more sophisticated alternative for XML, but also more complex

The End

- Listen to Captain Sparrow – this advice will come in handy during your degree.

