

Programming Design and Implementation

Lecture 7: Input / Output

Updated: 3rd May, 2020

Mark Upston

Discipline of Computing

School of Electrical Engineering, Computing and Mathematical Sciences (EECMS)

Copyright © 2020, Curtin University

CRICOS Provide Code: 00301J

Outline

Exceptions

File IO

Reading Files

File Handling

Parsing

Writing Files

“Zero” Marks

- ▶ A student who does any of the following in a submitted, assessable answer will receive heavy penalties, up to and including **zero** marks for that question:
 - ▶ Uses **continue**
 - ▶ Uses **break** in any other place than a **switch** statement
 - ▶ Uses **goto**
 - ▶ Has more than one **return** statement in a method
 - ▶ Has a **return** statement in a method anywhere but the last statement of the method
 - ▶ Uses **System.exit()** anywhere but the last statement of the **main()** method
 - ▶ Uses global variables for anything other than class fields
 - ▶ Uses a ternary operator
- ▶ Note: similar efforts in pseudo code will also receive zero marks

Exceptions - Reminder

- ▶ Error handling is a necessary task, but how do you do it elegantly?
 - ▶ Errors aren't "normal", you don't make a system that *expects* errors
 - ▶ But you **must** handle error situations
 - ▶ One solution: return an error code (see UCP COMP1000)
- ▶ Object-Oriented languages (such as Java) solve error handling with **exceptions**
 - ▶ An independent "return path" designed specifically for notifying the caller of an exceptional situation (error)
 - ▶ On an error, a method "throws" an exception
 - ▶ The calling method can "catch" the exception
 - ▶ If the caller doesn't catch it, the exception is thrown to the next highest caller
 - ▶ If no one catches it, the exception causes the program to crash

Exceptions - Reminder (2)

- ▶ Java only lets objects of type **Exception** or its descendants to be thrown
 - ▶ Java has a range of classes descending (inheriting, extends) from **Exception**
 - ▶ e.g., **IllegalArgumentException** and **ArrayIndexOutOfBoundsException**
 - ▶ You may define your own **Exception** class, as long as it inherits from **Exception** (or one of its subclasses)
 - ▶ This will be covered in DSA (COMP1002)

Catching Exceptions

- ▶ Exceptions from different methods in different objects are often all caught at the one place in the calling method
- ▶ Somewhere close to **main**
 - ▶ Convenient: all error handling happens in one place
- ▶ Most languages use **try**, **catch()**, **finally** blocks
 - ▶ **try**: Define the set of statements whose exceptions will all be handled by the catch block associated with this try
 - ▶ **catch()**: Processing to do if an exception is thrown in the try
 - ▶ **finally**: Processing to always do, regardless of whether an exception occurs or not
 - ▶ Good for cleanup, e.g., closing files
 - ▶ This block is optional and executed *after* the try and catch blocks

Input With Exception Handling

```
public static double realInput(String prompt, double lower, double upper)
{
    double num;
    Scanner sc = new Scanner(System.in);
    String outputPrompt = prompt;
    do
    {
        try
        {
            System.out.print(outputPrompt + " between " + lower + " and " +
                               upper);
            num = sc.nextDouble();
        }
        catch(InputMismatchException e)
        {
            sc.next(); //clear the buffer of the problematic input
            num = lower - 1.0; //set num to invalid so we don't exit
        }
        outputPrompt = "ERROR: please enter a valid value \n" + prompt;
    } while ((num < lower) || (num > upper));
    return num;
}
```

Why Files?

- ▶ RAM (Random Access Memory) is volatile and private to an application, so it's not a good match for the following purposes:
 - ▶ Storing application data long-term (between runs)
 - ▶ Sharing information between applications
 - ▶ Reading in bulk data provided by a user
- ▶ In contrast, files stored on disk are (semi)permanent, can be shared between applications and can be manipulated by the user outside the application

File Input/Output

- ▶ Unlike RAM, files are effectively an input to and/or an output from the application
 - ▶ Hence the term File I/O, for Input/Output
- ▶ Three basic steps in File I/O: (these apply to any platform)
 - ▶ Open the file
 - ▶ Read data from a file and/or write data to a file
 - ▶ Close the file

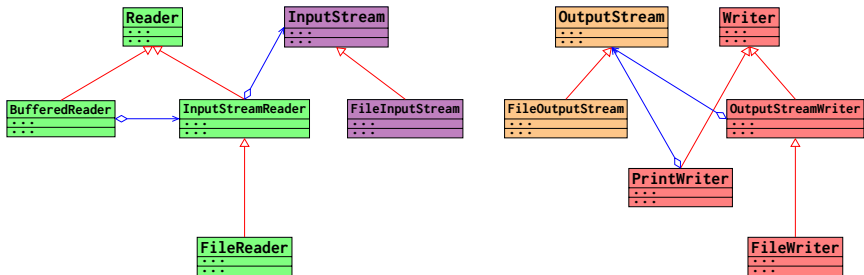
File I/O in Java

- ▶ Java encapsulates File I/O in a set of classes
 - ▶ All File I/O classes support the three basic steps
 - ▶ Different classes are used for manipulating different kinds of files or to perform I/O in different ways
 - ▶ Depends on what you need
 - ▶ Other kinds of I/O are also part of the same set of classes
 - ▶ e.g., Keyboard input, screen output, network I/O, device I/O, etc.
 - ▶ This was done to unify the handling of all the different I/O streams into a single consistent programming interface

Java I/O Classes

- ▶ The I/O classes can be split into three groups:
 - ▶ Classes that represent an I/O stream (e.g., file, device, etc.)
 - ▶ Classes to read from a stream
 - ▶ Classes to write to a stream
- ▶ There are specialised versions of the basic classes for each I/O type (file, network, device, etc.)
- ▶ We will concentrate on the file I/O classes
 - ▶ I am just making you aware of the broader scope to help explain why there are so many classes involved

Java I/O Class Hierachy



- **FileReader** and **FileWriter** are more for convenience and we won't be talking about them, or using them

Steps in Java for Reading a File

- ▶ 1. Create a stream object for a file
 - ▶ **FileInputStream**
- ▶ 2. Create an object that can read that stream
 - ▶ **InputStreamReader**
- ▶ 3. Read and process data from a file
- ▶ 4. Close the **FileInputStream**
- ▶ Note: A **FileReader** combines the first two steps into one

Steps in Java for Writing a File

- ▶ Basically the same as reading
- ▶ 1. Create a stream object for a file
 - ▶ **FileOutputStream**
- ▶ 2. Create an object that can write to that stream
 - ▶ **OutputStreamWriter** or **PrintWriter**
- ▶ 3. Write data to the file
- ▶ 4. Close the **FileOutputStream**
- ▶ Note: Again, **FileWriter** combines the first two steps into one

Efficient File I/O

- ▶ If you look at the methods for **InputStreamReader**, you'll see that it has two **read()** methods
 - ▶ One for reading a **byte** at a time
 - ▶ This returns **-1** if no more bytes to read (i.e., end-of-file reached)
 - ▶ A second for reading an array of **byte**'s at a time
- ▶ Since it's rare that we ever want just a single **byte** of data, the first method seems less useful
- ▶ However, when trying to read in text data there are conceptual difficulties with reading chunks at a time


Text vs. Binary Data

- ▶ Data files generally fall into two broad categories:
 - ▶ Files containing **text** data (relatively unstructured)
 - ▶ Files containing **binary** data (everything else)
- ▶ Binary data is highly structured
 - ▶ e.g., images, databases, executable files, etc
 - ▶ There is a lot of prior knowledge on precisely where information exists in the file and how large blocks are
- ▶ Text data is unstructured, so it is usually impossible to know beforehand how many bytes to read
 - ▶ e.g., how many characters in an arbitrary line of text?

Efficient File I/O Processing

- ▶ So reading blocks of N **byte**'s from a file is okay
 - ▶ You will know beforehand how many bytes are needed
- ▶ But when reading text, you have to constantly be ready for the end of data since you can't predict it
 - ▶ e.g., end of word, end of line, or end of file
- ▶ One approach is to read data in a **byte** at a time, and check each byte for end-of-X (' ', '\n' or -1)
 - ▶ The problem is that this will be very slow since hard disks are fastest at reading blocks of data at a time
 - ▶ Its like filling a bucket one drop at a time vs. opening the tap and letting the water flow freely into the bucket

InputStreamReader




- ▶ Note: **read()** returns **-1** when end-of-file (i.e., no more data) occurs
 - ▶ Also, note the '**\n**' (newline) is returned; just like any other character in the file

BufferedReader

- ▶ **BufferedReader** reads in chunks of data, buffering it in memory to search for end-of-X
 - ▶ In particular, the **readLine()** method looks for the end-of-line
 - ▶ The idea is to read in ('buffer') chunk of 1024 bytes, then search these bytes in RAM one at a time for the '**\n**'
 - ▶ '**\n**' = newline character, marking the end of the line in text
 - ▶ Since RAM is much faster than disk, this is much more efficient
 - ▶ When a line is found it is extracted and returned, but the rest of the buffer is kept in memory
 - ▶ Then for subsequent **readLine()** calls, it first checks the buffer for the '**\n**' before reading another 1024-byte chunk from the file
- ▶ Note: you create a **BufferedReader** from another **Reader**

BufferedReader (2)



- Note: **readLine()** does not include the '**\n**' in the returned lines and returns **null** when end-of-file occurs

File Reading: One Line at a Time - Pseudocode

```
SUBMODULE: readFile
IMPORT: filename (String)
EXPORT: none
  theFile := OPENFILE filename
  lineNum := 0
  INPUT line FROM theFile

  WHILE NOT (theFile = EOF) // EOF = end of file.
    // Detecting this is language-specific
    lineNum := lineNum + 1
    processLine(line) // Whatever line processing is required

    INPUT line FROM theFile // Read the next line
  ENDWHILE

  CLOSEFILE theFile // Close the file
END readFile
```

File Reading: One Line at a Time - Java

```
public void readFile(String inFileName)
{
    FileInputStream fileStream = null;
    InputStreamReader rdr;
    BufferedReader bufRdr;
    int lineNum;
    String line;
    try
    {
        fileStream = new FileInputStream(inFileName);
        rdr = new InputStreamReader(fileStream);
        bufRdr = new BufferedReader(rdr);
        lineNum = 0;
        line = bufRdr.readLine();
        while(line != null)
        {
            lineNum++;
            processLine(line);
            line = bufRdr.readLine();
        }
        ...
    }
}
```

File Reading: One Line at a Time - Java (2)

```
...

    fileStream.close();
}
catch(IOException e)
{
    if(fileStream != null)
    {
        try
        {
            fileStream.close();
        }
        catch(IOException ex2)
        { }
    }
    System.out.println("Error in fileProcessing: " +
                        e.getMessage());
}
}
```

Notes on File Handling

- ▶ Make sure you close the file as soon as possible
 - ▶ The operating system must track what files are open
 - ▶ The OS remembers where you were in the file, etc.
 - ▶ The resources available for this tracking are limited
 - ▶ Run out and the OS will terminate your program
 - ▶ Thus don't leave files open – clean them up early
 - ▶ Java doesn't free objects immediately - it waits for the garbage collector, so always explicitly **close()** a file once finished with it

File Handling - Exceptions

- ▶ You must handle **IOException**
 - ▶ **IOException**'s must be caught
 - ▶ Checked **Exception**
 - ▶ The compiler will complain otherwise
 - ▶ This also forces us to do the `try{} catch(){}` around **close**
 - ▶ Note: that you can add a `throws IOException` clause to the method signature (method contract) to avoid having to catch **IOException**'s
 - ▶ It just means that now the calling method must catch them.

Multiple catch Clauses

- ▶ Previous example caught all IO Exceptions
- ▶ What if we want to handle various exceptions in a different way?
 - ▶ e.g., if a file doesn't exist, ask the user for a different file name, but if anything else goes wrong with the file, terminate the program.
- ▶ Every **try** clause must be followed by one or more **catch** clauses
 - ▶ Order is important, it will attempt to catch the first **Exception** written first, if that is a less specific type then it may not catch a more precise type later
 - ▶ This concept will be explained more when you cover Inheritance in DSA (COMP1002)

File Reading: One Line at a Time - Java

```
public void readFileMultipleException() throws IOException
{ // Why would it be bad to write "throws Exception" here?
    Scanner sc = new Scanner(System.in);
    String inFileName; // File Code omitted - see slide 22
    boolean noFile;
    do
    {
        noFile = true;
        inFileName = sc.next();
        try
        {
            while(line != null)
            {
                processLine(line);
            }
            noFile = false;
        }
        catch(FileNotFoundException e)
        {
            System.out.println("Couldn't find your file " + e.getMessage() + " Try again!");
        }
        catch(IOException e)
        {
            throw e;
        }
        catch(Exception e)
        {
            throw new Exception("I am a bad programmer: ", e);
        }
    } while(noFile);
}
```

Parsing

- ▶ When dealing with text, it's often necessary to take it apart and organise it ready for processing/storage
- ▶ This is called **parsing** - to determine and extract the structure of a piece of text
 - ▶ The word originally comes from syntax analysis of written language
- ▶ Examples of where parsing is needed:
 - ▶ Natural language processing (e.g., spelling/grammar checks)
 - ▶ Building Web search indexes
 - ▶ Compilers - must parse the code to detect statements and variables

Tokenizing

- ▶ **Tokenizing** is the first step in parsing, the process of breaking up a stream of text into basic elements
 - ▶ We will use the US spelling (with a 'z') to avoid confusion
- ▶ These elements are called tokens, and what they are depends on what the application is parsing
 - ▶ e.g., Single words, entire lines, equation terms, etc.
- ▶ Tokens are broken up by searching for character(s) that delimit the boundary of a token
 - ▶ e.g., lines are separated by a '\n' newline character
 - ▶ e.g., words are separated by spaces, commas and periods
 - ▶ e.g., equation operands are separated by operators + / * -

Tokenizing with Java

- ▶ Java provides two classes to assist in tokenizing:
 - ▶ **StringTokenizer**
 - ▶ **StreamTokenizer**
 - ▶ They both do a similar job, except one works on **String**'s while the other works on **Readers**
 - ▶ For this unit we will focus on **split(String regex)** (below)
- ▶ Java also provides the **split(String regex)** method in the **String** class for simple tokenizing of a **String**
 - ▶ It is easier to use than the tokenizers, but won't return the delimiting character (which is something to remember)
 - ▶ We also don't know what regex is yet, but for the purpose of this unit, we can use basic regex to split on

Comma Separated Values

- ▶ We'll take a little detour and introduce comma separated values
 - ▶ We'll use this as an example for **split()** later on
- ▶ We often need to store data to a file
 - ▶ The question is, in what form should we store it?
- ▶ If the data is in table or matrix form, one can write it out as a set of rows and columns in a certain format:
 - ▶ One row is written per line
 - ▶ Each row contains multiple fields, one per column
 - ▶ Each field's value is separated by a comma **,**

CSV Example

► Contents of the file:

Title	Label Series1	Label Series2	Label Series3
Sales per region	Africa	Asia	Europe
Jan	34	67	56
Feb	36	87	78
Mar	31	56	88
Apr	29	67	92
Mar	43	56	78
May	54	71	68
Jun	42	65	82
X-values	Values Series 1	Values Series 2	Values Series 3

► The CSV file:

```
Sales per region,Africa,Asia,Europe
Jan,34,67,56
Feb,36,87,78
Mar,31,56,88
Apr,29,67,92
Mar,43,56,78
May,54,71,68
Jun,42,65,82
```


CSV Notes

- ▶ The delimiting commas have no trailing space
 - ▶ It may look better to have space after the comma, but it just makes it harder to parse when reading the file later!
- ▶ Column sizes don't have to be consistent across rows
 - ▶ e.g., the first row (headings) has much longer fields than the same columns in subsequent rows
 - ▶ It would be a waste of space to pad out the fields
 - ▶ Parsing/tokenizing must handle these variable length fields
- ▶ Numeric data is converted into its textual equivalent
 - ▶ If we saved integers directly, we might get things that look like text `'\n'` (ASCII 13) or `'.'` (ASCII 44) but are merely part of the data

Text vs. Binary

- ▶ You don't have to save table data in CSV format
 - ▶ In fact, dumping raw binary data is often more efficient, and you know how large each field is (e.g., ints = 4 bytes)
- ▶ However, CSV text data has some advantages:
 - ▶ Easy for humans to read and edit
 - ▶ Highly portable to different platforms
 - ▶ Big endian, vs. little endian
 - ▶ Fields are explicitly separated with commas
 - ▶ CSV is a widely-known format
 - ▶ XML has surpassed CSV as the standard format for data interchange between companies, but CSV is still used.

String.split() and CSV Data

- ▶ So lets parse a CSV file with String.split()
- ▶ The idea is you call a String variable's split method with what you want to "split" it on.
- ▶ It will return a **String[]** containing each part of the split in each element of the array
- ▶ Then we can iterate over the array to get the "parts" that we need out.

Parsing a Single CSV Row

```
private void processLine(String csvRow)
{
    String[] splitLine;
    splitLine = csvRow.split(",");

    for(int ii = 0; ii < splitLine.length; ii++)
    {
        System.out.print(splitLine[ii] + " ");
    }
    System.out.println("");
}
```

Output

- ▶ Given the following CSV data file:

97452,James,88,96,82,86

99576,Alan,6,46,34,38

9888,Geoff,100,68,72,75

- ▶ The following would be the output of having **readFileExample()** call the method **processLine()**:

97452 James 88 96 82 86

99576 Alan 6 46 34 38

9888 Geoff 100 68 72 75

Writing Text Files

- ▶ Writing files is actually simpler than reading them
 - ▶ Since you don't have to worry about parsing
- ▶ The overall approach is the same:
 - ▶ Open a **FileOutputStream**
 - ▶ Create a **Writer**
 - ▶ Output data to the file using the **Writer**
- ▶ One thing to be careful with is that you must ensure the newlines and commas are put in the right place
 - ▶ Assuming you are outputting CSV format of course

PrintWriter

- ▶ You can use **OutputStreamWriter**, but it is a bit clunky
 - ▶ The **write()** method requires an array of bytes so you have to copy your data into an array before writing
- ▶ It would be easier if you could write to a file in the same way you can print messages to the terminal
 - ▶ Fortunately you can use a **PrintWriter**
 - ▶ In fact, **System.out** is an instance of a **PrintWriter**
 - ▶ Thus writing to files can be identical to printing to the screen
 - ▶ Just don't forget to **close()** the file after you are done

PrintWriter Example

- ▶ The next slide shows an example of writing to a file
- ▶ It assumes you have passed in the student marks info that was read in earlier from the earlier CSV file
 - ▶ e.g., we are saving one line of the data to a new CSV file
 - ▶ Note that the commas are inserted between the fields, and there are no spaces added.

Writing a CSV Row

```
private void writeOneRow(String filename, int ID, String name, double assign,
                        double test, double exam, double overall)
{
    FileOutputStream fileStrm = null;
    PrintWriter pw;
    try
    {
        fileStrm = new FileOutputStream(filename);
        pw = new PrintWriter(fileStrm);
        pw.println(id + "," + name + "," + assign + "," + test + "," + exam + "," + overall);
        pw.close();
    }
    catch(IOException e)
    {
        if (fileStrm != null)
        {
            try
            {
                fileStrm.close();
            }
            catch(IOException ex2)
            { }
        }
        System.out.println("Error in writing to file: " + e.getMessage());
    }
}
```

Things Can be Simpler

- ▶ We have explored the 'long' way to open files for I/O
- ▶ Java provides other classes and constructors that do some of the steps for you
 - ▶ e.g., **FileReader**, **FileWriter**
 - ▶ e.g., the **FileOutputStream** creation can be done by an alternate constructor in **PrintWriter** that accepts the `fileName`
- ▶ I didn't show you the convenience classes / constructors so as to explicitly show you how Java is put together
 - ▶ But you must use the way we have shown you in your programs

Next Week

- ▶ The next Lecture will address the following:
 - ▶ Basic Sorting Algorithms